

Findings of the First GROUNDLM Shared Tasks: Evaluating Grounded Language Models across Visual and Scientific Evidence

Yimu Wang¹, Mozhgan Nasr Azadani¹, Xuye Liu¹, Yee Man Choi¹
Bo Xue², Krzysztof Czarnecki¹

¹University of Waterloo, ²City University of Hong Kong

Abstract

The first GROUNDLM shared tasks evaluate whether language models can identify the evidence behind an answer, rather than merely produce a correct answer in isolation. GOLD-ENVIEWVQA requires a model to answer a multiple-choice question about six synchronized driving-camera views and identify the supporting view. LITTRACEQA requires systems to retrieve scientific papers, locate answer-bearing evidence, and return multiple-choice or structured-table answers. We summarize the task design, submitted systems, official results, and recurrent failure modes. Among teams that submitted system reports, the best official joint accuracy on GOLDENVIEWVQA is 0.8814. Answer accuracy is consistently higher than view accuracy, while macro view accuracy remains lower than micro accuracy, exposing both attribution errors and view imbalance. On the required LITTRACEQA test set, the best score among teams that submitted system reports is 0.8002. Paper retrieval approaches saturation for several systems, but evidence F1 and structured-table accuracy remain substantially lower. Across both tasks, the central finding is consistent: producing a correct answer is easier than recovering the exact evidence representation required to justify it. We also report an optional 4,901-question LITTRACEQA diagnostic split, whose unverified labels make it useful for stress testing but not for ranking systems on the main task.

1 Introduction

Language models can often produce a plausible or correct answer without using the evidence a task intends them to use. In retrieval-augmented question answering, the answer may come from parametric memory rather than the cited document. In multi-view visual reasoning, the model may select the correct option while pointing to a camera that does not show the decisive object or relation. Answer-only evaluation hides both failures.

The first GROUNDLM shared tasks make evidence provenance part of the prediction. GOLD-ENVIEWVQA asks systems to jointly answer a driving-scene question and identify the camera view that supports the answer (Wang et al., 2026b). LITTRACEQA asks systems to retrieve relevant papers, identify supporting locations within those papers, and answer in a requested format (Liu et al., 2026). The tasks differ in modality and scale, but share a common structure: a prediction is correct only to the extent that its answer and its evidence trace are both correct.

This report describes the shared-task data and metrics, summarizes the submitted system reports, and analyzes official hidden-test results. We emphasize three findings. First, answer correctness and evidence identification remain separable: answer metrics are strong even when view or document-level grounding is weaker. Second, the dominant grounding bottleneck depends on the domain. GOLDENVIEWVQA systems struggle with view imbalance, cross-view identity, and the distinction between object visibility and decisive evidence. LITTRACEQA systems retrieve papers reliably but lose accuracy when mapping evidence to exact pages, object types, row keys, and table cells. Third, small development sets and repeated leaderboard interaction can make apparently strong improvements split-specific. Several reports explicitly document overfitting, post-hoc correction, or submission-level diagnosis, making evaluation protocol part of the scientific finding rather than a peripheral implementation detail.

2 Shared Task Setup

2.1 GoldenViewVQA

GOLDENVIEWVQA contains multi-view visual questions grounded in nuScenes driving scenes. Each instance provides a question, four answer options, and paths to six synchronized camera images:

front, front-left, front-right, back, back-left, and back-right. A system predicts an answer label and one supporting view; NONE_OF_THE_ABOVE is also a valid view label. Images are obtained separately under the nuScenes terms and are not redistributed with the benchmark.

Data curation. The source benchmark was constructed around conflict-centric nuScenes events rather than generic object-recognition questions (Wang et al., 2026b). Candidate events were mined from kinematic, geometric, and map features; the relevant actors and temporal windows were then identified, and their 3D boxes were projected into the six camera planes to propose candidate golden views. Questions target causality, counterfactual reasoning, and intent prediction. Annotators inspected all synchronized views to verify the answer, confirm that one view contained the most direct evidence, and revise or remove ambiguous cases; unsupported cases were retained under the no-support label. For the shared-task release, question identifiers and normalized question text are disjoint across development and test, and near-duplicate questions are screened, while scene and frame overlap is allowed so that complementary questions may probe the same context.

The shared-task release contains 55 labeled development examples and 59 input-only test examples. The development distribution is imbalanced: 38 of 55 examples use the front camera as the gold view (Mandiga, 2026; Gaafar and Saqr, 2026). This makes both micro and class-balanced view evaluation necessary.

Metrics. Answer and view accuracy are exact-match rates over the predicted answer option and view label. View macro accuracy is the unweighted mean of recall over gold view classes, giving each represented class equal weight despite the skew toward the front camera. Joint accuracy is $N^{-1} \sum_i \mathbf{1}[\hat{a}_i = a_i^* \wedge \hat{v}_i = v_i^*]$: an example receives credit only when both predictions are correct. Joint accuracy is the primary leaderboard score; the three component metrics are reported for diagnosis.

2.2 LitTraceQA

LITTRACEQA evaluates literature-grounded question answering over a metadata pool of 27,487 papers. Given a question and an answer contract, a system returns canonical paper identifiers, coarse evidence locators, and an answer. Evidence may be a text span, table, figure, equation or algorithm,

or citation context. Answers may be free-form text, a multiple-choice label, or a table conforming to a supplied schema (Liu et al., 2026).

Data curation. LITTRACEQA uses a generate-ground-challenge pipeline over recent machine-learning, computer-vision, and NLP papers (Liu et al., 2026). An open-book generator receives paper text and a target evidence category, then proposes a question together with gold papers, typed evidence items, and an answer. Automated checks require every evidence item to belong to a declared gold paper and require distinctive supporting text or values to be recoverable from the source. Closed-book challenger models then attempt the question without paper access, filtering candidates that can be answered reliably from model memory. The required hidden test set was subsequently screened by human auditors; the much larger test_extra split did not receive that final manual screening.

The public validation set contains 55 questions: 26 hidden-source single-paper questions and 29 multi-paper questions. The required hidden test set contains 71 questions, comprising 50 multiple-choice and 21 table-answer questions; no free-form question is scored on this split. The optional test_extra split contains 4,901 automatically collected questions. Its labels are hidden but not human-verified, evidence is not evaluated, and participation does not affect the main leaderboard. We therefore treat it only as a diagnostic stress test.

Metrics. Paper retrieval is compared as a set of canonical paper identifiers: precision, recall, and F1 are computed per question and then macro-averaged. Evidence uses the same set-based metrics after each locator is normalized to the tuple (paper_id, source_type, page_or_section, object_id); fine-grained regions and table cells are not part of this evidence key. Multiple-choice answers use accuracy and free-form answers use exact match. For tables, rows are aligned by the schema’s row-key columns: row F1 evaluates the aligned row keys, macro cell accuracy first computes accuracy per question, and micro cell accuracy pools all non-key cells. On the required test set, the answer aggregate is $A = (\text{MCAcc} + \text{RowF1} + \text{CellAcc}_{\text{macro}})/3$, and the leaderboard score is $(\text{PaperF1} + \text{EvidenceF1} + A)/3$. On test_extra, evidence is not scored, so the score is $(\text{PaperF1} + A)/2$. This shared-task composite differs from the strict joint diagnostic proposed by Liu et al. (2026).

2.3 Result Reporting Policy

Official results were obtained from the [GROUNDLM evaluation server](#). The main tables in this paper include teams that submitted a system report to the workshop and are sorted by their best leaderboard score. They are therefore a report-linked view of the leaderboard, not a count of all registered teams or submissions. A team’s best server submission does not always correspond to the particular configuration emphasized in its paper. We preserve this distinction when interpreting methods and results.

3 Submitted Systems

3.1 GoldenViewVQA Systems

GroundSight. CoVer-VQA starts from a zero-shot joint prediction and progressively applies view-specific verification, confidence-aware joint verification, and group-level checking across semantically related questions ([Wang et al., 2026a](#)). The answer is frozen during its view-specific stage, limiting collateral changes to already correct answers. The report’s four-stage pipeline reaches 0.8475 joint accuracy, while the final 0.8814 submission includes two additional evaluator-informed post-hoc corrections that the authors explicitly separate from the general pipeline.

Maddies. The Maddies report contrasts a decomposition-heavy cascade with Joint Answer-Preserving Visual Re-grounding (JAVR) ([Channa et al., 2026](#)). JAVR anchors on a strong joint prediction, freezes the answer, localizes the queried target across views, revisits disputed cameras with crops, and changes only the view when a strict visual certificate is satisfied. This restrained controller changes few predictions and is motivated by the observation that repeated verification can replace correct outputs rather than improve them.

Ensar. Ensar detects agents in all six views, back-projects detections to the ground plane, associates objects across cameras, and serializes the scene as text with explicit per-object view sets ([Gaafar and Saqr, 2026](#)). A language model receives either the serialization alone or the serialization together with the images. Development experiments show that pixels substantially improve view attribution while leaving answer accuracy nearly unchanged. The report also identifies failures caused by serializing missed detections as confirmed negative evidence.

aravind. The aravind report provides a deliberately non-visual baseline: always predict CAM_FRONT and use Qwen3-4B to answer from the question, answer options, and non-visual metadata ([Mandiga, 2026](#)). The reported run obtains 0.7119 joint accuracy without accessing images, although the team’s best leaderboard entry is 0.7288. This baseline isolates the extent to which view priors and textual cues can support performance without visual grounding.

3.2 LitTraceQA Systems

DKE. SAGE converts each query into typed evidence obligations, retrieves papers through several routes, parses PDFs into provenance-preserving records, and selects a bounded package using slot-aware multimodal reranking and coverage priorities ([Xu et al., 2026](#)). Benchmark-facing identifiers and locators are projected only from retrieved records. Relative to its clean baseline, the reported system raises hidden-test paper F1 from 0.7199 to 0.9915 and evidence F1 from 0.3179 to 0.7643.

gabby. TRACE preserves target identity and cardinality across retrieval, evidence localization, and schema-constrained extraction ([Gupta and Godara, 2026](#)). It treats a paper as a source container rather than assuming one paper equals one output row, and validates all output against the evaluator contract. The report highlights a clean-track artifact scoring 0.7606. The team’s best server score, 0.7923, corresponds to an adaptive artifact that the authors discuss diagnostically rather than claim as the transferable automatic system.

tus-nlp. Lit3R combines question-anchored retrieval with paper-to-paper expansion before local reading and cross-paper synthesis ([Ise et al., 2026](#)). Its Retrieve stage iteratively decomposes questions and verifies evidence sufficiency; its Relate stage expands from likely seed papers using semantic similarity, bibliographic coupling, and lexical similarity. On validation, adding this expansion raises multi-paper recall@50 from 0.647 to 0.940.

OdeD. OdeD begins with a conventional dense-retrieval, reranking, and multimodal-reading pipeline, then studies the evaluator and submission history as measurement instruments ([Saqr and Gaafar, 2026](#)). Its best fully automated score is 0.5519; the final 0.7649 score includes per-question PDF auditing and changes attributed through aggregate leaderboard feedback. The report discloses this

Team	Joint	View	V-Macro	Answer
GroundSight	.8814	.8983	.7464	.9661
Maddies	.7966	.8136	.5930	.9492
Ensar	.7797	.8305	.6522	.9153
aravind	.7288	.7627	.5580	.9322

Table 1: Best official GOLDENVIEWVQA results for teams with a submitted system report. V-Macro is macro view accuracy. The table is sorted by joint accuracy and does not enumerate teams without a report.

adaptation and argues that its final score should not be read as expected performance on unseen questions.

ThousandOfWind. ThousandOfWind combines word and character TF-IDF, BGE-M3 retrieval, query-derived source planning, PDF page and object ranking, deterministic structural extractors, and guarded Qwen fallbacks (Zhirui Zhu, 2026). Its development process retains only type-level rules that survive multi-document probes and label-isolated execution. The paper emphasizes that narrow rules can raise development scores while reducing the credibility of the resulting system.

Everest. PageRecall combines sparse and dense paper retrieval, whole-document reading when context permits, structural anchors for numbered objects, permutation voting for multiple choice, and a reporting-only entailment audit (Chauhan, 2026). A controlled development analysis finds that BM25 top-4 page selection exposes a gold page only 52.6% of the time, whereas the reader cites the correct page in 45 of 48 emitted locators when the gold page is available. This motivates removing page selection for papers that fit in context.

4 Official Results

Table 1 shows that all four report-linked submissions answer more accurately than they identify the supporting view. The gap between answer and view accuracy ranges from 6.78 points for GroundSight to 16.95 points for aravind; the gap between answer and joint accuracy is larger still. View macro accuracy is 15.19–22.06 points below view accuracy, showing that aggregate view correctness remains concentrated in frequent classes. GroundSight obtains the strongest result on every reported metric, but its paper distinguishes the general four-stage system from the two evaluator-informed changes included in the final submission (Wang et al., 2026a).

The LITTRACEQA results in table 2 expose a

clear pipeline gradient. DKE and tus-nlp obtain 0.9915 paper F1, and four teams reach perfect multiple-choice accuracy. Evidence F1 is lower for every team, ranging from 0.4407 to 0.7643. Structured tables remain harder: row F1 peaks at 0.5927, and macro cell accuracy peaks at 0.4040. Thus retrieving the right source and selecting a multiple-choice answer do not imply that a system can recover the exact evidence tuple or bind values to the correct table structure.

Table 3 reports the optional large split separately. ThousandOfWind retains non-zero performance across retrieval and answer metrics, while the tus-nlp submission receives zero paper F1 and near-zero table scores. These results show that the split can expose severe distribution or pipeline failures, but the absence of human verification, evidence labels, and broad submission coverage prevents fine-grained system ranking. The split is most useful as a scale stress test and a source of candidates for future manual annotation.

5 Findings and Discussion

5.1 Correct Answers Do Not Guarantee Correct Evidence

The strongest cross-task pattern is the separation between answering and attribution. On GOLDENVIEWVQA, answer accuracy exceeds view accuracy for every report-linked team, and non-visual systems already obtain high answer accuracy. Ensar reports 0.9455 development answer accuracy from its text serialization without raw images, while pixels primarily improve the view head (Gaffar and Saqr, 2026). The aravind baseline reaches 0.8475 test answer accuracy with no visual input at all (Mandiga, 2026). Even the visually grounded Maddies system ends with nine test examples whose answer is correct but whose view is wrong (Channa et al., 2026). These results explain why joint evaluation is necessary: answer accuracy alone substantially overstates evidence-grounded performance.

The same separation appears in LITTRACEQA. SAGE identifies two multiple-choice examples with correct answers but no true-positive submitted evidence (Xu et al., 2026). PageRecall’s entailment audit similarly finds that whether an answer is supported by its cited page is nearly independent of whether it is correct on its development data (Chauhan, 2026). On the official test set, multiple-choice accuracy is 0.92–1.00 for all report-linked

Team	Score	Paper F1	Evid. F1	MC Acc.	Row F1	Cell-M	Cell- μ
DKE	.8002	.9915	.7643	1.0000	.5927	.3413	.3793
gabby	.7923	.9869	.7277	1.0000	.5825	.4040	.4483
tus-nlp	.7803	.9915	.7369	1.0000	.5288	.3087	.3218
OdeD	.7649	.9704	.7319	.9800	.4992	.2976	.3563
ThousandOfWind	.7374	.9869	.6620	1.0000	.4198	.2698	.2759
Everest	.5768	.7618	.4407	.9200	.4214	.2421	.2874

Table 2: Best official results on the required 71-question LITTRACEQA test set for teams with a submitted system report. Evid. is evidence; Cell-M and Cell- μ are macro and micro cell accuracy. Free-form exact match is not applicable on this split. The table reports best team submissions, which may differ from the clean or fully automated configurations emphasized in individual reports.

Team	Score	P-F1	MC	Row	Cell-M
ThousandOfWind	.5267	.7390	.5215	.3216	.1001
tus-nlp	.0257	.0000	.1501	.0032	.0010

Table 3: Diagnostic LITTRACEQA test_extra results for report-submitting teams with a leaderboard entry. The 4,901-question split is optional, its labels are not human-verified, and evidence is not evaluated; these scores are not comparable to the main leaderboard.

teams, whereas evidence F1 is at most 0.7643. A model may answer from parametric knowledge, a nearby passage, or a plausible option without satisfying the explicit provenance contract.

5.2 Grounding Fails at Interfaces, Not Only in Base Models

For GOLDENVIEWVQA, the hard distinction is often not whether an object is visible, but which view establishes the relation asked about. CoVeR-VQA reports errors in which a salient object attracts the model to a close-up camera even though another view better shows its position relative to the lane and ego vehicle (Wang et al., 2026a). Mad-dies reaches a similar design conclusion from the opposite direction: a large cascade can overwrite correct decisions, so its final controller freezes the answer and grants only narrow authority to change the view (Channa et al., 2026). Ensar shows that perception-to-text interfaces introduce their own failures: a missed detection becomes an apparently certain negative, or cross-camera association assigns an object to the wrong view set (Gaafar and Saqr, 2026). Better base models alone do not remove these representation errors.

For LITTRACEQA, the bottleneck shifts along a longer chain. Lit3R demonstrates that paper-to-paper expansion can raise multi-paper recall@50 from 0.647 to 0.940 (Ise et al., 2026), and the leaderboard confirms that paper retrieval is close to

saturated for several systems. The remaining errors concern identity at finer granularity. PageRecall finds page selection to be an upstream bottleneck even when the correct paper has already been found (Chauhan, 2026). TRACE describes a “grounding contract gap” in which correct source access still fails because a locator, observation unit, canonical row key, or value-to-source association is wrong (Gupta and Godara, 2026). SAGE likewise finds exact row-key and cell-value composition harder than recovering table structure (Xu et al., 2026). OdeD’s dominant error is the right paper but the wrong row or column, not an invented number (Saqr and Gaafar, 2026). These findings favor typed intermediate representations and provenance-preserving interfaces over undifferentiated end-to-end generation.

5.3 Shortcuts and Imbalance Matter

GOLDENVIEWVQA exposes two shortcuts. First, the development view distribution makes an always-front predictor strong in micro accuracy but weak in macro accuracy. The aravind baseline records 0.7797 test view accuracy from this constant rule, while its macro view accuracy is 0.1667 for the run analyzed in the paper (Mandiga, 2026). Second, answer-option wording carries substantial non-visual signal. Both aravind and Ensar report high answer accuracy without images (Mandiga, 2026; Gaafar and Saqr, 2026). Macro view accuracy partially corrects the first shortcut, but the primary joint metric remains influenced by class frequency. Future editions should preserve macro reporting, screen answer options with language-only baselines, and increase coverage of minority and unsupported view classes.

LITTRACEQA has different evaluator-facing shortcuts. Evidence credit depends on an exact tuple of paper, source type, page, and object identifier; semantically correct adjacent evidence can therefore score zero. Structured-table met-

rics additionally depend on matching row keys before cell values receive credit. OdeD shows that these rules make the score arithmetically diagnosable from controlled submissions (Saqr and Gaafar, 2026). TRACE and ThousandOfWind respond with contract-constrained normalization and label-isolated audits rather than unrestricted generation (Gupta and Godara, 2026; Zhirui Zhu, 2026). These mechanisms improve scorer compatibility, but the task should continue to distinguish genuine grounding from reproduction of annotation conventions.

5.4 Small Development Sets Complicate Selection

Both tasks provide 55 labeled development examples. Several reports show that this is enough to guide design but too small for stable model or heuristic selection. Ensar observes a reversal between two models: the run with higher development macro view accuracy generalizes substantially worse on test (Gaafar and Saqr, 2026). Maddies finds that increasing reasoning effort or adding verification stages does not reliably improve the cascade (Channa et al., 2026). CoVeR-VQA obtains large gains from staged verification, but its final two corrections rely on evaluator feedback and are therefore reported separately (Wang et al., 2026a).

The LITTRACEQA reports are unusually explicit about the same issue. ThousandOfWind removes high-scoring development heuristics that encode named benchmarks, methods, or expected answer patterns (Zhirui Zhu, 2026). PageRecall shows that padding predictions to match development-set paper cardinality transfers poorly and reduces paper F1 by 0.245 in an intermediate test configuration (Chauhan, 2026). OdeD and TRACE distinguish reusable system behavior from adaptive artifacts constructed through repeated aggregate-score diagnosis (Saqr and Gaafar, 2026; Gupta and Godara, 2026). Consequently, the best leaderboard score should not automatically be interpreted as the best estimate of out-of-sample system performance. Future shared tasks would benefit from a separate final-submission phase, lower feedback granularity near the deadline, or an additional untouched evaluation split.

6 Conclusion

The first GROUNDLM shared tasks evaluate a capability that answer-only benchmarks miss: tracing an answer to the visual or scientific evidence that supports it. The submitted systems achieve strong answer and paper-retrieval performance, yet leave substantial headroom in view identification, exact evidence localization, and structured extraction. Across modalities, the main errors occur at interfaces—between objects and views, papers and pages, or source tables and canonical rows—and can remain invisible when only the final answer is scored. The results support continued joint evaluation of answers and evidence, class-balanced grounding metrics, stronger non-visual and retrieval-only baselines, and evaluation protocols that separate reusable system improvements from adaptation to a small benchmark.

Limitations

This findings report relies on official aggregate scores and the analyses supplied in system reports. The main tables intentionally include only teams with submitted reports, so they are not a complete enumeration of leaderboard activity. Best leaderboard submissions may differ from the clean, automated, or reproducible configuration emphasized by a report, and we call out known cases rather than attributing every best score to the named method. The hidden test labels prevent independent per-example error analysis. Both main test sets are small, and some GOLDENVIEWVQA classes contain very few examples. Finally, LITTRACEQA test_extra is not human-verified and lacks evidence evaluation, so its results should not be used to draw fine-grained conclusions about grounding quality.

References

- Aditya Channa, Madhav Kapoor, and Jastej Singh. 2026. Maddies at GroundLM 2026 shared tasks: Answer-preserving visual re-grounding for multi-view driving VQA. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Aaditya Chauhan. 2026. PageRecall: Measuring page selection in literature-grounded question answering. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.

- Mostafa Gaafar and Omar Saqr. 2026. Ensar at GroundLM 2026 Shared Tasks: Symbolic Scene Serialisation for View-Level Evidence Identification. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Sachin Gupta and Divya Godara. 2026. TRACE: Target-aware retrieval, attributed evidence, and contract-constrained extraction for LitTraceQA. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Akira Ise, Kotaro Kumagai, Yuta Yamaguchi, Hisanori Ozaki, Yukio Uematsu, and Ikuya Yamada. 2026. Lit3R: Retrieve-relate-read for evidence-grounded question answering over scientific literature. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Xuye Liu, Yimu Wang, Peng Shi, Bo Xue, Xiangrui Ke, Songcheng Cai, Kath Choi, Di Wu, Freda Shi, and Krzysztof Czarnecki. 2026. [Littraceqa: A benchmark for multi-stage grounding and verification in scientific question answering](#).
- Aravind Mandiga. 2026. A non-visual baseline for multi-view visual question answering on golden-viewvqa. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Omar Saqr and Mostafa Gaafar. 2026. Oded at groundlm 2026 shared tasks: Reading the scorer for literature-grounded qa. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Kun Wang, Yupeng Hu, Ruping Cao, Hao Liu, Zhiran Li, Qianlong Xiang, and Harry Cheng. 2026a. Covervqa: Multi-stage evidence verification for golden-viewvqa. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Yimu Wang, Yee Man Choi, Barry Zhang, Mozgan Nasr Azadani, Sean Sedwards, and Krzysztof Czarnecki. 2026b. [Where does the answer come from? benchmarking view-level visual evidence identification in multi-view mllms for autonomous driving](#).
- Gang Xu, Ce Huang, Yushan Pan, and Wei Wang. 2026. Sage: Slot-aware grounded evidence acquisition for littraceqa. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.
- Zhirui Zhu. 2026. ThousandOfWind at GroundLM 2026 shared task: An auditable hybrid pipeline for LitTraceQA. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*.